

# Assignment 11

Deadline: Friday **02.02.2024** 23:55

## Notes:

- Solve the assignment **on your own** — no groups allowed.
- Hand-written solutions will not be accepted, except for graphs and diagrams.
- If you hand in non-pdf files or multiple files, name your submission as `stla23_11_SURNAME.zip`, replacing `SURNAME` with your surname. Otherwise `stla23_11_SURNAME.pdf`. Also include your full name in the submitted PDF.
- Submit your solution via **Ilias**.

This is a bonus assignment. It's not mandatory to solve it but you can get up to 20 bonus points if you do.

## Exercise 11.1 Vector Clocks

**20 Points**

In a distributed system, it is often necessary to determine a ordering of events that happened on the different processes. Because it is not possible to determine a global time in an asynchronous distributed system, we can only determine a partial ordering of events. One way to do this is by using *vector clocks*. A vector clock can be used to determine a happened before relation between events denoted by  $A \rightarrow B$  where  $A$  and  $B$  are events. For an event  $A$  on process  $i$  and an event  $B$  on process  $j$ ,  $A \rightarrow B$  holds if and only if

- $i = j$  and  $A$  happened before  $B$  on process  $i$ , or
- $A$  is the sending of a message on process  $i$  and  $B$  is the receiving of the same message on process  $j$  or
- there is an event  $C$  such that  $A \rightarrow C$  and  $C \rightarrow B$ .

Events  $A$  and  $B$  are said to be concurrent if neither  $A \rightarrow B$  nor  $B \rightarrow A$  holds.

A vector clock is a mapping from processes to natural numbers. Each process has its own vector clock. Let  $VC_i$  denote the vector clock of process  $i$ . The initial value of  $VC_i$  is  $VC_i[j] = 0$  for all processes  $j$ . When an event  $A$  happens on process  $i$ , the entry of  $VC_i$  for process  $i$  is incremented by one, i.e.,  $VC'_i[i] = VC_i[i] + 1$ . When process  $i$  sends a message to process  $j$ , it increments the entry of  $VC_i$  for process  $i$  by one and sends the vector clock  $VC_i$  with the message. When process  $j$  receives the message, it increments the entry of  $VC_j$  for process  $j$  by one and updates the entry of  $VC_j$  for all processes  $k$  by taking the maximum of the current value and the value received in the message, i.e.,  $VC'_j[k] = \max(VC_j[k], VC_i[k])$ . The sending and receiving of a message is also considered an event. An example of this is shown in Figure 1.

For an event  $A$  we denote the vector clock of the process on which  $A$  happened by  $VC(A)$ . Let  $VC(A)$  and  $VC(B)$  be two vector clocks. We state that  $VC(A) < VC(B)$  if and only if




Figure 1: Vector clock example with three processes taken from <https://sookocheff.com/post/time/vector-clocks/>

$VC(A)[i] \leq VC(B)[i]$  for all processes  $i$  and  $VC(A)[j] < VC(B)[j]$  for at least one process  $j$ . The following holds:

$$A \rightarrow B \Leftrightarrow VC(A) < VC(B)$$

With this we can determine a partial ordering of events.

- a) Write the module **VectorClock** in  $TLA^+$  that specifies the vector clock operations for  $N$  processes. Keep in mind that in asynchronous system there can happen events between the sending of a message and the receiving of the same message. Use constants to make the state space of the specification finite. **(10 points)**
- b) Write the specification **VectorClockMC** in  $TLA^+$  that keeps track of the happened before relation between events. Use the relation to check with the help of TLC if your specification **VectorClock** does imply the the happened before relation. **(10 points)**

---

**Total: 20 Points**